

Last Time:

- Regularization + Line Search w/ constraints
- Deterministic Optimal Control Problem

Today:

- Pontryagin's Minimum Principle
- Linear-Quadratic Regulator Intro

* Pontryagin's Minimum Principle

- Also called "maximum principle" if you maximize a reward
- First-order necessary conditions for deterministic optimal control problem
- In discrete time, just a special case of KKT
- Given:

$$\min_{\substack{x_{1:N} \\ u_{1:N}}} \sum_{n=1}^{N-1} l(x_n, u_n) + l_F(x_N)$$

$$\text{s.t. } x_{n+1} = f(x_n, u_n)$$

- We can form the Lagrangian:

$$L = \sum_{n=1}^{N-1} l(x_n, u_n) + \lambda_{n+1}^+ (f(x_n, u_n) - x_{n+1}) + l_F(x_N)$$

- This result is usually stated in terms of the "Hamiltonian"

$$H(x, u, \lambda) = l(x, u) + \lambda^T f(x, u)$$

- Plug it into L:

$$L = H(x_1, u_1, \lambda_1) + \left[\sum_{n=2}^{N-1} H(x_n, u_n, \lambda_{n+1}) - \lambda_n^T x_n \right] + l_F(x_N) - \lambda_N^T x_N$$

note change to indexing

- Take derivatives w.r.t. x and λ :

$$\frac{\partial L}{\partial \lambda_n} = \frac{\partial H}{\partial \lambda_n} - x_{n+1} = f(x_n, u_n) - x_{n+1} = 0$$

$$\frac{\partial L}{\partial x_n} = \frac{\partial H}{\partial x_n} - \lambda_n^T = \frac{\partial l}{\partial x_n} + \lambda_{n+1}^T \frac{\partial f}{\partial x_n} - \lambda_n^T = 0$$

$$\frac{\partial L}{\partial x_N} = \frac{\partial l_F}{\partial x_N} - \lambda_N^T = 0$$

- For u we write the min explicitly to handle torque limits:

$$u_n = \underset{\tilde{u}}{\operatorname{argmin}} \quad H(x_n, \tilde{u}, \lambda_{n+1})$$

$$\text{s.t. } \tilde{u} \in \mathcal{U}$$

shorthand for "in feasible set"

$$\text{e.g. } u_{\min} \leq \tilde{u} \leq u_{\max}$$

- Fn Summary:

$$x_{n+1} = \nabla_\lambda H(x_n, u_n, \lambda_{n+1}) = f(x_n, u_n)$$

$$\lambda_n = \nabla_x H(x_n, u_n, \lambda_{n+1}) = \nabla_x l(x_n, u_n) + \left(\frac{\partial f}{\partial x}\right)^T \lambda_{n+1}$$

$$u_n = \underset{\tilde{u}}{\operatorname{argmin}} \quad H(x_n, \tilde{u}, \lambda_{n+1})$$

s.t. $\tilde{u} \in \mathcal{U}$

$$\lambda_n = \frac{\partial l_F}{\partial x_n}$$

- These can be stated almost identically in continuous time:

$$\dot{x} = \nabla_\lambda H(x, u, \lambda) = f(x, u) \quad \begin{matrix} \text{continuous time} \\ \downarrow \text{dynamics} \end{matrix}$$

$$\dot{\lambda} = \nabla_x H(x, u, \lambda) = \nabla_x l(x, u) + \left(\frac{\partial f}{\partial x}\right)^T \lambda$$

$$u = \underset{\tilde{u}}{\operatorname{argmin}} \quad H(x, \tilde{u}, \lambda)$$

s.t. $\tilde{u} \in \mathcal{U}$

$$\lambda_n = \frac{\partial l_F}{\partial x_n}$$

* Some Notes:

- Historically, many algorithms were based on integrating the continuous ODEs for $X(t)$, $\lambda(t)$ forward/backward to do gradient descent on $U(t)$
- These are called "indirect" and/or "shooting" methods.
- In continuous time $\lambda(t)$ is called "co-state" trajectory
- These methods have largely fallen out of favor as computers + solvers have improved

* Linear-Quadratic Regulator (LQR)

$$\min_{\substack{X \in \mathbb{R}^{n \times n} \\ U \in \mathbb{R}^{n \times m}}} \underbrace{\sum_{n=1}^{N-1} \frac{1}{2} x_n^T Q_n x_n + \frac{1}{2} u_n^T R_n u_n}_{Q \geq 0, R > 0} + \frac{1}{2} x_N^T Q_N x_N$$

s.t. $x_{n+1} = A_n x_n + B_n u_n$

- Goal is to drive the system to the origin
- Called "time-invariant" LQR if $A_n = A$, $B_n = B$
 $Q_n = Q$, $R_n = R \quad \forall k$
- Called "time-varying" (TVLQR) otherwise.
- Typically time-invariant LQR is for stabilizing an equilibrium and TVLQR is for tracking trajectories.

- Can (locally) approximate many nonlinear problems
 - Many extensions e.g. infinite-horizon, stochastic
 - Has been called "crown jewel of control theory"

* LQR with Indirect Shooting

$$X_{n+1} = Ax_n + Bu_n$$

$$\lambda_n = Qx_n + A^T\lambda_{n+1}, \quad \lambda_n = Qx_n$$

$$u_n = -R^{-1}B^T \lambda_{n+1} \quad \leftarrow \text{this gives the gradient of the cost w.r.t. } u_n$$

-Procedure:

- 1) Start with an initial guess trajectory
 - 2) Simulate ("rollout") to get $X(t)$
 - 3) Backward pass to get $\lambda^{(t)}$ and $\Delta u^{(t)}$
 - 4) Rollout with line search on Δu
 - 5) Go To 3 until convergence

* Example :

- "Double Integrator": $\dot{x} = \begin{bmatrix} \dot{x} \\ \ddot{x} \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} q \\ \dot{q} \end{bmatrix} + \begin{bmatrix} 0 \\ 1 \end{bmatrix} u$

- Think of this as a "sliding brick" on ice (no friction)

$$X_{n+1} = \underbrace{\begin{bmatrix} 1 & h \\ 0 & 1 \end{bmatrix}}_A \underbrace{\begin{bmatrix} q_n \\ \dot{q}_n \end{bmatrix}}_U + \underbrace{\begin{bmatrix} \frac{1}{2}h^2 \\ h \end{bmatrix}}_B u_n$$

* Bryson's Rule

- Cost-tuning heuristic
- Set diagonal elements of Q and R to $(1/\max \text{ value})^2$
- Normalizes all cost terms to 1
- Good starting point for tuning.
- Scalar case:

$$\underbrace{\frac{1}{2} Q x^2}_{\left(\frac{1}{x_{\max}}\right)^2} + \underbrace{\frac{1}{2} R u^2}_{x_{\max}^2} = \left(\frac{1}{u_{\max}}\right)^2 u_{\max}^2 = 1$$