

Last Time:

- Dynamic Programming
- Convex Optimization

Today:

- Convex MPC
-

* Convex MPC

- Think of this as "constrained LQR"
- Remember from DP, if we have the cost-to-go we can get U_k by solving a one-step problem:
 - $U_k = \underset{u}{\operatorname{argmin}} \quad l(x_k, u) + V_{k+1}(f(x_k, u))$
 - $= \underset{u}{\operatorname{argmin}} \frac{1}{2} u^T R_u u + (A_k x_k + B_k u)^T P_{k+1} (A_k x_k + B_k u)$
- We can add constraints to this one-step problem, but performance will be bad because $V_h(x)$ was computed without constraints.
- There's no reason I can't add more steps to the one-step problem:

$$\min_{\substack{x_{1:H} \\ u_{1:H-1}}} \sum_{n=1}^{H-1} \frac{1}{2} x_n^T Q_n x_n + \frac{1}{2} u_n^T R_n u_n + x_H^T P_H x_H$$

$$\text{s.t. } x_{n+1} = A_n x_n + B_n u_n$$

$x_n \in \mathcal{X}$, \mathcal{X} , \mathcal{U} convex

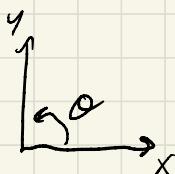
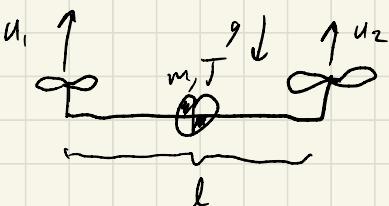
$$u_n \in \mathcal{U}$$

(e.g. $U_{\min} \leq u_n \leq U_{\max}$)

- $H \leq N$ is called the "horizon"
- With no additional constraints MPC / "receding horizon" exactly matches LQR for any H
- Intuition: explicit constrained optimization over first H steps gets the state close enough to the origin/reference so that constraints are no longer active and LQR solution is valid further into the future.
- In general:
 - A good approximation of $V(x)$ is important for good performance
 - Better $V(x) \Rightarrow$ shorter horizon will work
 - Longer $H \Rightarrow$ better performance, less reliant on quality of $V(x)$ approximation

* Example

- Planar Quadrotor



$$m\ddot{x} = -(u_1 + u_2) \sin(\theta)$$

$$m\ddot{y} = (u_1 + u_2) \cos(\theta) - mg$$

$$J\ddot{\theta} = \frac{l}{2}(u_2 - u_1)$$

- Linearize about hover:

$$\Rightarrow u_1 = u_2 = \frac{1}{2}mg \quad (\text{cancel gravity})$$

$$\Rightarrow \begin{cases} \Delta \ddot{x} \approx -g\theta \\ \Delta \ddot{y} \approx \frac{1}{m}(\Delta u_1 + \Delta u_2) \\ \Delta \ddot{\theta} \approx \frac{1}{J} \frac{1}{2}(\Delta u_2 - \Delta u_1) \end{cases}$$

$$\begin{bmatrix} \Delta \dot{x} \\ \Delta \dot{y} \\ \Delta \dot{\theta} \\ \Delta \ddot{x} \\ \Delta \ddot{y} \\ \Delta \ddot{\theta} \end{bmatrix} = \underbrace{\begin{bmatrix} 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}}_A \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \\ \Delta \dot{x} \\ \Delta \dot{y} \\ \Delta \dot{\theta} \end{bmatrix} + \underbrace{\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}}_B \begin{bmatrix} \Delta u_1 \\ \Delta u_2 \end{bmatrix}$$

- MPC Cost Function:

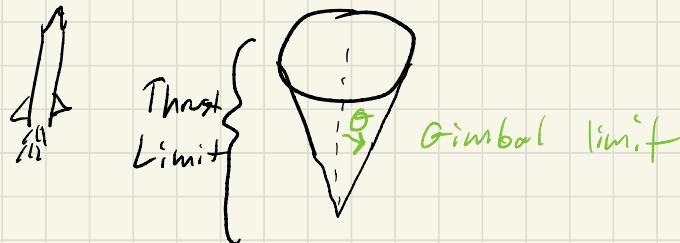
$$J = \sum_{n=1}^H \frac{1}{2} (x_n - x_{ref})^T Q (x_n - x_{ref}) + \frac{1}{2} \Delta u_n^T R \Delta u_n + (x_H - x_{ref})^T P_H (x_H - x_{ref})$$

$$\underbrace{\frac{1}{2} x_n^T Q x_n - \underbrace{x_{ref}^T Q x}_{q^T} + \dots}_{\text{...}}$$

* Other MPC Examples

- Rocket Landing

- Thrust vector constraints are naturally expressed as a cone:

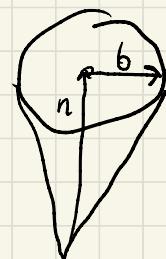


- Mars powered descent and SpaceX use SOCP-based MPC with linearized dynamics
- Legged Locomotion
- Contact forces must obey "friction cone" constraint

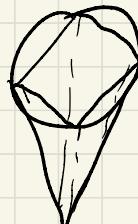
$$\|b\|_2 \leq \mu n$$

Friction Force Friction Coefficient normal force

Diagram illustrating the friction cone constraint. It shows a cone with a horizontal axis labeled b and a vertical axis labeled n . A bracket labeled "Friction Force" points to the horizontal axis, and a bracket labeled "Friction Coefficient" points to the vertical axis.



- Cone is often approximated as a pyramid so the constraint is linear ($Ax \leq b$)



- MIT Cheetah and other quadrupeds use QP-based MPC with linearized body dynamics + friction pyramid

• What about nonlinear dynamics?

- Linear stuff often works well, so use it if you can.
- Nonlinear dynamics are risky because optimization problems are non-convex
- No convergence guarantees
- Can work well in practice due to warm starting
(solution doesn't change much over iteration)
- Active area of research
- Common in autonomous driving