

Last Time:

- Convex MPC

Today:

- Nonlinear Trajectory Optimization
- Differential Dynamic Programming a.k.a.
Iterative LQR

* Nonlinear Traj Opt Problem

$$\min_{\substack{x_1 \\ u_1 \\ \dots \\ x_N \\ u_N}} J = \sum_{n=1}^{N-1} l_n(x_n, u_n) + l_N(x_N)$$

nonlinear/nonconvex

s.t. $x_{n+1} = f(x_n, u_n)$ ← dynamics constraints
(nonlinear)

$$x_n \in X_n \quad \left. \begin{array}{l} \\ \end{array} \right\}$$

Other state/input constraints
(nonlinear/nonconvex)

$$u_n \in U_n$$

- Usually we assume that costs are C^2
(continuous 2nd derivatives)
- Constraints are usually assumed to be C^1
(continuous first derivatives)

* Differential Dynamic Programming (DDP)

- Nonlinear traj opt method based on approximate DP
- Use 2nd order Taylor approximations of cost-to-go in DP recursion to compute improvements to our input trajectory and local feedback policy
- Can think of this as a mash-up of Newton's method and DP
- Very fast convergence is possible

- Cost-to-go expansion:

$$V_n(x + \Delta x) \approx V_n(x) + \underbrace{p_n^T \Delta x}_{\text{gradient}} + \frac{1}{2} \Delta x^T P_n \Delta x \quad \begin{matrix} \text{gradient} \\ \text{Hessian} \end{matrix}$$
$$p_n = \nabla_x l_n(x) \quad P_n = \nabla_{xx}^2 l_n(x)$$

- Action-Value function expansion:

$$S_n(x + \Delta x, u + \Delta u) \approx S_n(x, u) + \begin{bmatrix} g_x \\ g_u \end{bmatrix}^T \begin{bmatrix} \Delta x \\ \Delta u \end{bmatrix} + \frac{1}{2} \begin{bmatrix} \Delta x \\ \Delta u \end{bmatrix}^T \begin{bmatrix} G_{xx} & G_{xu} \\ G_{ux} & G_{uu} \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta u \end{bmatrix}$$
$$(G_{xu} = G_{ux}^T)$$

- Bellman Equation:

$$V_{n-1}(x + \Delta x) = \min_{\Delta u} \left[S(x, u) + g_x^T \Delta x + g_u^T \Delta u + \frac{1}{2} \Delta x^T G_{xx} \Delta x + \frac{1}{2} \Delta u^T G_{uu} \Delta u + \frac{1}{2} \Delta x^T G_{xu} \Delta u + \frac{1}{2} \Delta u^T G_{ux} \Delta x \right]$$

$$\frac{\partial []}{\partial \Delta u} = g_u + G_{uv} \Delta u + G_{ux} \Delta x = 0$$

$$\Rightarrow \Delta u_{n-1} = -G_{uv}^{-1} g_u - G_{uv}^{-1} G_{ux} \Delta x$$

$$\begin{aligned} & \text{"feed-forward term"} \\ & \Delta u_{n-1} = -d_{n-1} - K_{n-1} \Delta x \quad \text{"feedback term"} \end{aligned}$$

- Plug Δu back into S_{n-1} to recover $V_{n-1}(x + \Delta x)$

$$\begin{aligned} \Rightarrow V_{n-1}(x + \Delta x) & \approx V_{n-1}(x) + g_x^T \Delta x + g_u^T (-d_{n-1} - K_{n-1} \Delta x) \\ & + \frac{1}{2} \Delta x^T G_{xx} \Delta x + \frac{1}{2} (d_{n-1} + K_{n-1} \Delta x)^T G_{uv} (d_{n-1} + K_{n-1} \Delta x) \\ & - \frac{1}{2} \Delta x^T G_{xu} (d_{n-1} + K_{n-1} \Delta x) - \frac{1}{2} (d_{n-1} + K_{n-1} \Delta x)^T G_{ux} \Delta x \end{aligned}$$



$$P_{n-1} = G_{xx} + K_{n-1}^T G_{uv} K_{n-1} - G_{xu} K_{n-1} - K_{n-1}^T G_{ux}$$

$$p_{n-1} = g_x - K_{n-1}^T g_u + K_{n-1}^T G_{uv} d_{n-1} - G_{xu} d_{n-1}$$

- Need some more math to calculate g and G ...

* Matrix Calculus

- given $f(x) : \mathbb{R}^n \rightarrow \mathbb{R}^m$, let's look at 2nd-order Taylor expansion:

- If $m=1$:

$$f(x+\Delta x) \approx f(x) + \underbrace{\frac{\partial f}{\partial x} \Delta x}_{\mathbb{R}^{m \times n}} + \frac{1}{2} \underbrace{\Delta x^\top \frac{\partial^2 f}{\partial x^2} \Delta x}_{\mathbb{R}^{n \times n}}$$

- If $m > 1$:

$$f(x+\Delta x) \approx f(x) + \underbrace{\frac{\partial f}{\partial x} \Delta x}_{\mathbb{R}^{m \times n}} + \frac{1}{2} \underbrace{\left(\frac{\partial^2 f}{\partial x} [\frac{\partial f}{\partial x} \Delta x] \right) \Delta x}_{\text{ugly!}}$$

- for $m > 1$, $\frac{\partial^2 f}{\partial x^2}$ is a 3rd-rank tensor. Can think of this like a "3D matrix". We need more notation to keep track of which dimensions we're multiplying along.

- Kronecker Product:

$$\underbrace{A}_{l \times m} \otimes \underbrace{B}_{n \times p} = \begin{bmatrix} a_{11} B & a_{12} B & \dots \\ a_{21} B & a_{22} B & \ddots \\ \vdots & & \end{bmatrix}$$

$l \times m p$

(multiply each element of A by matrix B)

- Vectorization Operator

$$A = [A_1 \ A_2 \ \dots \ A_m]$$

$l \times m$ q
 column vectors

$$\text{vec}(A) = \begin{bmatrix} A_1 \\ A_2 \\ \vdots \\ A_m \end{bmatrix}$$

$l \times m \times 1$

- The "vec trick"

$$\text{vec}(ABC) = (C^T \otimes A) \text{vec}(B)$$

$$\Rightarrow \text{vec}(AB) = (B^T \otimes I) \text{vec}(A) = (I \otimes A) \text{vec}(B)$$

- If we want to diff a matrix w.r.t. a vector, vectorize the matrix:

$$\frac{\partial A_{mn}}{\partial x} = \underbrace{\frac{\partial \text{vec}(A)}{\partial x}}_{l \times n} \quad (\text{implied whenever we diff a matrix})$$

- Back to the Taylor expansion of $f(x)$:

$$f(x + \Delta x) \approx f(x) + \underbrace{\frac{\partial f}{\partial x} \Delta x}_{\mathbf{A}} + \frac{1}{2} (\Delta x^T \otimes I) \underbrace{\frac{\partial^2 f}{\partial x^2} \Delta x}_{\frac{\partial \text{vec}(\frac{\partial f}{\partial x})}{\partial x} \Delta x}$$

$\frac{\partial f}{\partial x}$ Δx $\frac{\partial^2 f}{\partial x^2}$

$\text{vec}(I \Delta x)$ $\frac{\partial \text{vec}(\frac{\partial f}{\partial x})}{\partial x}$

$(\Delta x^T \otimes I) \text{vec}(A)$ $\frac{\partial \Delta x}{\partial x}$

- Sometimes we need to diff through a transpose:

$$\frac{\partial}{\partial x} (A \alpha^T B) = (B^T \otimes I) T \underbrace{\frac{\partial A}{\partial x}}_{\text{"commutator matrix"}}$$

"commutator matrix"

$$T \text{vec}(A) = \text{vec}(A^T)$$

- Action - Value Function Derivatives:

$$S_n(x, u) = l_n(x, u) + V_{n+1}(f(x, u))$$

$$\Rightarrow \frac{\partial S}{\partial x} = \frac{\partial l}{\partial x} + \underbrace{\frac{\partial V}{\partial f} \frac{\partial f}{\partial x}}_A \Rightarrow \boxed{g_x = \nabla_x l(x, u) + A^T \nabla V(f(x, u)) \\ = \nabla_x l + A^T p_{n+1}}$$

$$\frac{\partial S}{\partial u} = \frac{\partial l}{\partial u} + \underbrace{\frac{\partial V}{\partial f} \frac{\partial f}{\partial u}}_B \Rightarrow \boxed{g_u = \nabla_u l(x, u) + B_n^T \nabla V(f(x, u)) \\ = \nabla_u l + B_n^T p_{n+1}}$$

$$\boxed{G_{xx} = \frac{\partial g_x}{\partial x} = \nabla_{xx}^2 l(x, u) + A_n^T \nabla^2 V_{n+1} A_n + (p_{n+1}^T \otimes I) T \frac{\partial A_n}{\partial x}}$$

$$G_{uu} = \frac{\partial g_u}{\partial u} = \nabla_{uu}^2 l(x, u) + B_n^T \nabla^2 V_{n+1} B_n + (p_{n+1}^T \otimes I) T \frac{\partial B_n}{\partial u}$$

$$G_{xu} = \frac{\partial g_x}{\partial u} = \nabla_{xu}^2 l(x, u) + A_n^T \nabla^2 V_{n+1} B_n + (p_{n+1}^T \otimes I) T \frac{\partial B_n}{\partial u}$$

$$G_{ux} = \frac{\partial g_u}{\partial x} = \nabla_{ux}^2 l(x, u) + B_n^T \nabla^2 V_{n+1} A_n + (p_{n+1}^T \otimes I) T \frac{\partial A_n}{\partial x}$$

↑ Tensor terms

- Often tensor terms are left out because they are expensive to compute. Equivalent to doing Gauss-Newton instead of full Newton.
- Version without tensor terms called : LQR
- Forward Rollout :

$$\Delta J = 0$$

for $K=1:N-1$

$$u'_n = u_n - \underbrace{\alpha}_{\text{line search step length}} d_n - K_n (x'_n - x_n) \quad \begin{matrix} \leftarrow \text{new value} \\ \leftarrow \text{old value} \end{matrix}$$

$$x'_{n+1} = f(x'_n, u'_n)$$

$$\Delta \tilde{J} \leftarrow \Delta J + \underbrace{\alpha g_{u_n}^T d_n}_{\substack{\text{expected cost change based on} \\ \text{1st-order Taylor expansion}}}$$

end

- Line Search:

$$\alpha = 1$$

do :

$$x', u', \Delta \tilde{J} = \text{rollout}(x, u, d, K, \alpha)$$

$$\alpha \leftarrow \underbrace{\zeta \alpha}_{\text{back tracking parameter } \sim 1/2}$$

$$\text{while } J(x', u') > J(x, u) - b \Delta \tilde{J}$$

\nwarrow Armijo tolerance
 $\sim 10^{-2}$

- Repeat until convergence

+ Example

- Acrobot swing-up
- DDP can converge in fewer iterations but iterations are more expensive. iLQR often wins in wall clock time.
- Problem is nonconvex \Rightarrow can land in different local optima depending on initial guess