

Last Time:

- Nonlinear trajectory optimization
- DDP / iLQR

Today:

- DDP details + extensions
- Constraints
- Free / minimum time problems

* DDP Recap:

- Solve the unconstrained traj opt problem:

$$\min_{\begin{array}{l} x_{1:N} \\ u_{1:N} \end{array}} \sum_{n=1}^{N-1} l(x_n, u_n) + l_N(x_N)$$

$$\text{s.t. } x_{n+1} = f(x_n, u_n)$$

- Backward Pass:

$$V_n(x+\Delta x) \approx V(x) + \frac{1}{2} \Delta x^T P_n \Delta x + p_n^T \Delta x$$

$$P_n = \nabla^2 l_n(x) \quad p_n = \nabla l_n(x)$$

$$V_{n-1}(x+\Delta x) = \min_u S(x+\Delta x, u+\Delta u)$$

$$\Rightarrow \Delta u_{n-1} = - \Delta u_n - K_{n-1} \Delta x_n$$

$$\begin{aligned} P_{n-1} &= G_{xx} + K^T f_w K - G_{xu} K - K^T G_{ux} \\ p_{n-1} &= g_x - K^T g_u + K^T G_{ud} - G_{xu} d \end{aligned}$$

- Forward Rollout:

$$\Delta J = 0$$
$$x'_1 = x_1$$

for $k = 1:N-1$

$$u'_n = u_n - \alpha d_n - K_n (x'_n - x_n)$$

$$x'_{n+1} = f(x'_n, u'_n)$$

$$\Delta J \leftarrow \Delta J + \alpha g_n^\top d_n$$

end

- Line Search:

$$\alpha = 1$$

do:

$$x', u', \Delta J = \text{rollout}(x, u, d, K, \alpha)$$

$$\alpha \leftarrow c \alpha$$

$$\text{while } J(x', u') > J(x, u) - b \Delta J$$

$$x, u \leftarrow x', u'$$

- Repeat until $\|d\|_\infty < \text{tol}$

- Regularization

- Just like standard Newton, Vars and/or Sosn Hessians can become indefinite in the backward pass
- Definitely necessary for DDP, often a good idea with iLQR as well.
- Many options for regularizing
- Add a multiple of identity to $D^2 \ell(x,u)$ just like standard Newton.
- Regularize P_h as needed in the backward
- Regularize $G_{uu} = D_{uu}^2 S(x,u)$ as needed in the backward pass (remember this is the only matrix we have to invert)

$$K = G_{uu}^{-1} G_{ux}, \quad d = G_{uu}^{-1} g_u$$

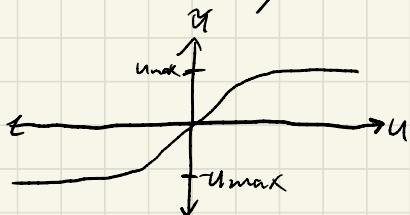
- This last one is a good choice for iLQR but not DDP.
- Regularization shouldn't be required with iLQR but can be necessary due to roundoff/truncation error that accumulates in the backward pass.

* DDP Notes:

- Can be very fast (iterations + wall clock time)
- One of the most efficient traj opt method since it exploits DP structure
- Always dynamically feasible due to forward rollout \Rightarrow can always execute on robot.
- Comes with TULQR tracking controller for free.
 - \Rightarrow Can be very effective for online use.
- Does not natively handle constraints
- Does not support infeasible initial guess for state trajectory due to forward rollout. Bad for "maze"/"bug trap" scenarios.
- Can suffer from numerical ill-conditioning, especially for long trajectories, due to error accumulation in the backward pass. Careful problem scaling can be needed to make it work.

* Handling Constraints with DDP

- Many options depending on type of constraint
- Torque limits are often handled heuristically with a "squashing function" e.g. $\tanh(u) = \frac{u_{\max} - u}{u_{\max} + u}$:



$$\tilde{u} = u_{\max} \tanh\left(\frac{u - u_{\min}}{u_{\max} - u_{\min}}\right)$$

- Effective, but adds nonlinearity and may need more iterations to converge.
- Better option: solve box-constrained QP in the backward pass!

$$\Delta u = \underset{\Delta u}{\operatorname{arg\,min}} \quad S(x_{k+1}, u + \Delta u)$$

$$\text{s.t. } u_{\min} \leq u + \Delta u \leq u_{\max}$$

- Also need to zero out rows of K corresponding to active torque limit constraints
- State constraints are harder. Often penalties are added to cost function. Can cause ill-conditioning.
- Better option: wrap entire DDP algorithm in an augmented Lagrangian method.
- AL method adds linear (multiplier) and quadratic (penalty) terms to cost function so easily fits into DDP.

- Very effective for achieving fast convergence to coarse constraint tolerances, but can fail to achieve high accuracy due to ill-conditioning

* Handling free initial state:

- Allow the solver to choose X_0 , possibly within some bounds.
- Add fictitious $k=0$ step to the problem with dynamics:

$$\begin{aligned} X_0 &= f(X_0, u_0) = u_0 \\ &= Ax_0 + Bu_0 \end{aligned}$$

A \nearrow x_0
 B \nearrow u_0

- Now u_0 becomes X_0

* Handling free/minimum time problems:

$$\begin{aligned} \min_{\substack{x(0) \\ u(t)}} J &= \int_0^T 1 dt \quad \underbrace{\qquad}_{\text{min time}} \\ T &\qquad \text{problem} \\ \text{s.t. } \dot{x} &= f(x, u) \\ x(T) &= x_{\text{goal}} \\ u_{\min} &\leq u(t) \leq u_{\max} \end{aligned}$$

- We don't want to change the number of knot points.
- Make h (time step) from RK a control input

$$x_{n+1} = f(x_n, u_n), \quad \tilde{u} = \begin{bmatrix} u \\ h \end{bmatrix}$$

- Also scale cost by h , e.g.:

$$J(x, \tilde{u}) = \sum_{n=1}^{N-1} h_n l(x_n, u_n) + l_N(x_N)$$

- This formulation is always nonlinear/non-convex even if dynamics are linear.
- Requires constraints on h , otherwise the solver can "cheat physics" by making h negative or very large to exploit discretization error in the RK method.