

Last Time:

- Hybrid methods for contact

Today:

- Reasoning about friction
- Iterative Learning Control

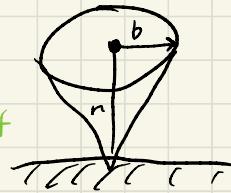
* Friction in Trajectory Optimization

- In our hybrid TrajOpt setup we assumed sticking (no slip).
- Equivalent to infinite friction
- We can enforce additional friction cone constraints for each contact point to make sure the robot won't slip:

$$\|b\|_2 \leq Mn$$

↑ ↓
friction force $\in \mathbb{R}^2$ friction coefficient

normal force $\in \mathbb{R}$



- This constraint is often linearized:

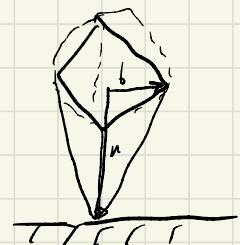
$$e^T d \leq Mn \quad , \quad d \in \mathbb{R}^n, \quad e = \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix}$$

$d \geq 0$

$$b = \begin{bmatrix} F & -F \end{bmatrix} d$$

2×2

$$= \sum \text{abs}(b_i) = \|b\|_1$$



- The "friction pyramid" constraint is easy to enforce in QP-based MPC to avoid foot slip.
 - Hybrid methods generally try to avoid. If you want to model stick-slip behavior, need to add more modes.
 - Friction is hard to model. Coulumb is only an approximation. Typically use a conservatively small value of M .
-

* What Happens when our model has errors?

- Models are always approximations
- Simpler models are often preferred even if they're less accurate.
- Feedback (e.g. LQR or MPC) can often compensate for model error.
- Sometimes that isn't good enough (e.g. very tight constraints, performance requirements)
- Several Options:

- 1) Parameter Estimation: Classical "grey-box" system identification. Fit e.g. masses in your model using data from the real system.
 - + Very simple efficient
 - + Generalizes
 - Assumes model structure
- ~ Improve the model

2) Learn Model: Fit a generic function approximator to full dynamics or residual.
Classical "black-box" system identification.

- + Doesn't assume model structure
- + Generalizes
- Not sample efficient. Requires lots of data.

3) Learn a Policy: Standard reinforcement learning approach: Optimize a function approximator for feedback policy.

- + Makes few assumptions
- Doesn't generalize
- Not sample efficient. Requires many "rollouts"

4) Improve a Trajectory: Assuming we have a reference trajectory computed with nominal model, Improve it with data from the real system.

- + Makes few assumptions
- Assumes you have a decent model
- Doesn't generalize
- + Very sample efficient

Improve the policy

* Iterative Learning Control

- Can think of this as a very specialized policy gradient method for the policy class

$$u_n = \bar{u}_n - K_n(x_n - \bar{x}_n)$$

y can be any tracking controller
reference control input

where we are improving the reference \bar{u}

- Can also think of this as SQP method where we get the "right-hand side" vector from a rollout on the actual system.
- Assume we have a reference trajectory \bar{x}, \bar{u} that we want to track:

$$\min_{\substack{x_1 \\ u_1 \\ \vdots \\ x_{N-1} \\ u_{N-1}}} J = \sum_{n=1}^{N-1} \frac{1}{2} (x_n - \bar{x}_n)^T Q (x_n - \bar{x}_n) + \frac{1}{2} (u_n - \bar{u}_n)^T R (u_n - \bar{u}_n)$$

(can be any cost function) + $\frac{1}{2} (x_N - \bar{x}_N)^T Q_N (x_N - \bar{x}_N)$

$$\text{s.t. } x_{n+1} = f(x_n, u_n)$$

- As we've seen before, the KKT system for this problem looks like:

$$\begin{bmatrix} H & C^T \\ C & 0 \end{bmatrix} \begin{bmatrix} \lambda \\ z \end{bmatrix} = \begin{bmatrix} -\nabla J \\ -C(z) \end{bmatrix}$$

$$\text{where } z = \begin{bmatrix} x_1 \\ u_1 \\ \vdots \\ x_N \end{bmatrix}, \quad C(z) = \begin{bmatrix} \vdots \\ f(x_n, u_n) - \bar{x}_{n+1} \\ \vdots \end{bmatrix}, \quad C = \frac{\partial C}{\partial z},$$

Gauss-Newton Hessian \rightarrow

$$H = \begin{bmatrix} Q & R \\ R^T & 0 \end{bmatrix}, \quad G = \begin{bmatrix} Q_N \\ G \end{bmatrix}$$

- Note that we have ∂J in the RHS instead of ∂L and λ instead of ∂J in the LHS. We can do this because the Lagrangian is linear in λ . We normally prefer solving for ∂J because that lets us regularize.
- Two important observations:
 - 1) If we do a rollout on the real system, $(\bar{z}) = 0$ always (for the true dynamics).
 - 2) Since we know J , given x_n, u_n from a rollout, we can compute ∂J
- Now we have the RHS vector for the KKT system. We also know b from the cost function.
- We can compute $C = \frac{\partial C}{\partial z}$ using x_n, u_n and the nominal model.
- However, since our nominal model is approximate and assuming x_n, u_n is already close to \bar{x}, \bar{u} , we can just use $C = \frac{\partial C}{\partial z} |_{\bar{x}, \bar{u}}$, which can be pre-computed once.
- Now just solve the KKT system for ∂z , update $\bar{u} \leftarrow \bar{u} + \Delta u$, and repeat.
- Can easily add inequality constraints (e.g. torque limits) and solve a QP.

* Why Should ILC Converge?

- We've already seen approximations in Newton's method: Gauss-Newton, regularization.
- In general, these are called "inexact" and "quasi-Newton" methods. Many variants (e.g. BFGS, Broyden), with well-developed convergence theory.
- For a generic root-finding problem:

$$f(x + \Delta x) \approx f(x) + \underbrace{\frac{\partial f}{\partial x} \Delta x}_\text{Exact Newton Step} = 0$$

- As long as Δx satisfies $\|f(x) + \frac{\partial f}{\partial x} \Delta x\| \leq \eta \|f(x)\|$ for some $\eta < 1$, an inexact Newton method will converge.
- Convergence is slower than exact Newton
- This means we can use $J \approx \frac{\partial f}{\partial x}$ to compute Δx

* ILC Algorithm:

do :

$\Delta x, \Delta u \leftarrow \underset{\Delta x, \Delta u}{\operatorname{argmin}} \quad \mathcal{T}(\Delta x, \Delta u)$

s.t. $\Delta x_{n+1} = A_n \Delta x_n + B_n \Delta u_n$

$u_{\min} \leq u_n \leq u_{\max}$

$\bar{u} \leftarrow \bar{u} + \Delta u$

while $\|x_n - \bar{x}_n\| \geq tol$

whole trajectory from real system

$x_n, u_n \leftarrow \text{rollout}(\bar{x}, \bar{u})$ (on real system)

This is a QP